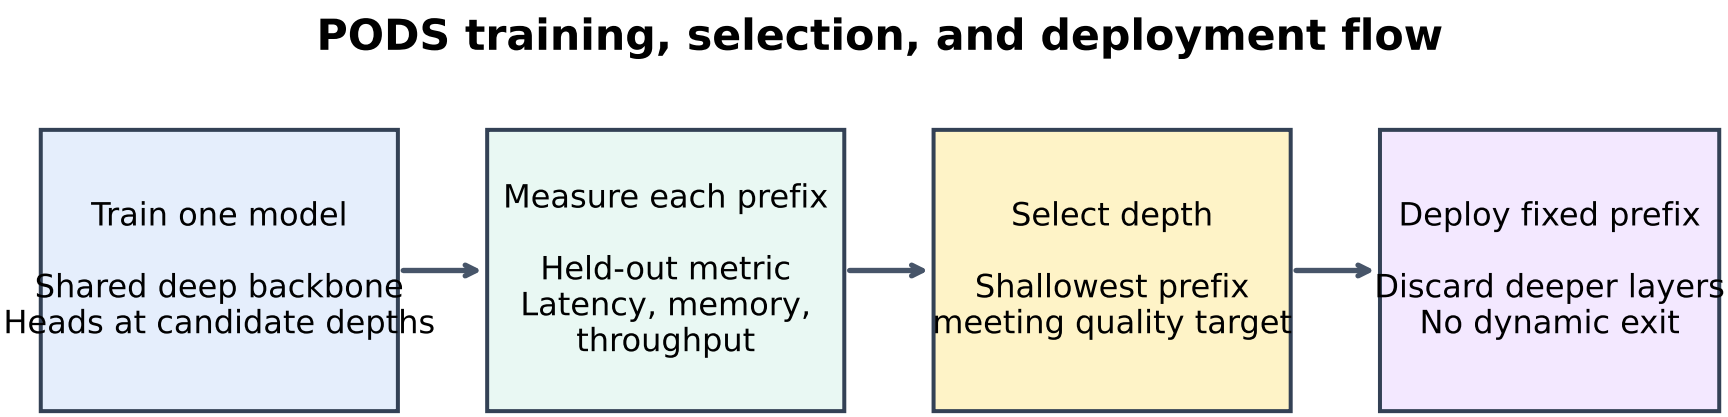
PODS training, selection, and deployment flow
Train one model
Shared deep backbone
Heads at candidate depths
Measure each prefix
Held-out metric
Latency, memory, throughput
Select depth
Shallowest prefix meeting quality target
Deploy fixed prefix
Discard deeper layers
No dynamic exit